

# Coprime-Factor Security Architecture — Blueprint v0.1

**Purpose.** Define a defense-in-depth model where each protective layer is engineered to be *mathematically independent* from the others—analogous to coprime moduli. The goal is to minimize common-mode failures and make successful compromise require simultaneous defeat of multiple, pairwise-independent layers.

---

## 1) Executive Summary

- **Problem:** Conventional defense-in-depth often hides correlated risks (same crypto family, same supply chain, same control plane). One bug or key leak can pierce multiple layers.
  - **Idea:** Engineer layers as **coprime factors**—each using disjoint primitives, trust roots, vendors, failure modes, and ops paths—so that compromise probabilities multiply instead of correlate.
  - **Outcome:** Lower breach probability, provable separation of concerns, auditable independence, graceful degradation.
- 

## 2) Principles & Independence Criteria

Design layers so that for any pair ( $L_i, L_j$ ):

- **Cryptographic independence:** different algorithm *families* (e.g., AES-GCM  $\leftrightarrow$  ChaCha20-Poly1305; Ed25519  $\leftrightarrow$  P-256; RSA  $\leftrightarrow$  ECDSA) and distinct RNG sources.
- **Trust root independence:** separate CAs/KMS/HSMs (ideally different vendors + firmware trees).
- **Codebase independence:** different libraries, compilers, languages; no shared critical deps (e.g., not both OpenSSL).
- **Runtime independence:** different isolation tech (VM  $\leftrightarrow$  container  $\leftrightarrow$  enclave/microVM), different kernels when feasible.
- **Ops path independence:** separate teams, approval flows, credentials, monitoring stacks.
- **Supply-chain independence:** distinct artifact pipelines, signing roots, registries.
- **Control-plane independence:** policy engines and orchestration stacks that cannot impersonate each other.

**Independence Rule:** If any two layers share a single catastrophic common dependency, they are *not* coprime.

---

## 3) Formal Model (Concise)

Let ( $L_1, \dots, L_k$ ) be layers; ( $C_i$ ) is event “Layer  $i$  compromised” with probability ( $p_i$ ). We target **approximate independence:** ( $P(C_i \wedge C_j) \approx P(C_i)P(C_j)$ ) for all ( $ij$ ). Then overall protection event ( $C = \bigvee_i C_i$ ) yields breach probability ( $P(C) \approx \sum_i (1 - p_i)$ ). Independence is enforced via the criteria above; audits measure residual correlation.

**Independence Score (IS):** For each dimension  $d \in \{\text{crypto, trust-root, codebase, runtime, ops, supply, control}\}$ , assign  $(w_d)$ . For layer pair  $(i,j)$ , set  $(s_d)^{\{(i,j)\}}$  (independent or not). Then  $[IS = \sum_{i < j} d w_d, s_d^{\{(i,j)\}}]$  Target **IS**  $\geq 0.85$  before production.

---

## 4) Reference Architecture (High Level)

**User  $\leftrightarrow$  Edge  $\leftrightarrow$  App Plane  $\leftrightarrow$  Data Plane  $\leftrightarrow$  Control Plane  $\leftrightarrow$  Audit Plane**

**Layer A — Identity & Access (Coprime AuthN):** - FIDO2/WebAuthn (ECDSA P-256) on vendor HSM-A - PLUS independent TOTP/HOTP (HMAC-SHA1/256) or passkey on HSM-B - Optional biometric check locally; never elevates trust by itself

**Layer B — Transport & Session:** - mTLS (TLS 1.3, library L1) *and* secondary WireGuard tunnel (library L2) for admin paths - Different cipher families (AES-GCM vs ChaCha20-Poly1305)

**Layer C — Authorization (Dual Policy Consensus):** - OPA/Rego engine and AWS Cedar (or equivalent) run in separate control planes - High-risk actions require **AND-consensus**; low-risk allow OR with rate-cap

**Layer D — Data At Rest (Dual Encryption):** - Envelope encryption with AES-GCM (KMS-A) - Nested layer with ChaCha20-Poly1305 (KMS-B) - Keys held in HSMs from different vendors with distinct firmware

**Layer E — Integrity & Build Trust:** - Provenance via Sigstore (Fulcio/Rekor) + separate in-house CA - Reproducible builds verified by two independent verifiers

**Layer F — Observability & Audit:** - Append-only Merkle log (e.g., Trillian) inside org - Daily external anchoring (public transparency log or blockchain) with separate signer

**Layer G — Runtime Isolation:** - Service set S1 on microVM/Firecracker; S2 on containers with SELinux/AppArmor; disjoint kernels when possible

---

## 5) Pattern Catalogue (Implementable)

**P1. Dual-Primitive Encryption (DPE):** - Apply two independent AEAD schemes with unrelated keys from independent KMS/HSM - Decrypt path requires both layers; failure leaves data safe

**P2. Dual Policy Consensus (DPC):** - Two policy engines evaluate requests; decision = AND for privileged ops - Engines have separate repos, CI, and deploy channels

**P3. Split-Vendor Key Ceremony (SVK):** - Root keys generated in two vendor HSMs, different RNGs; quorum requires both

**P4. Divergent Transport (DT):** - Administrative channels traverse a second, distinct VPN/overlay using disjoint crypto

**P5. Twin Attestation (TA):** - Workload must present attestation from two independent roots (e.g., TPM quote + SEV-SNP/TEE report)

---

## 6) Example Config (Pseudo-YAML)

```
authn:
  factors:
    - type: fido2
      curve: p256
      hsm: vendorA
    - type: totp
      algo: hmac-sha1
      hsm: vendorB
  require: AND

transport:
  primary_tls:
    lib: rustls
    aead: aes-256-gcm
  admin_overlay:
    type: wireguard
    aead: chacha20-poly1305

authorization:
  engines:
    - name: opa
      source_repo: git://corp/opa-policies
    - name: cedar
      source_repo: git://corp/cedar-policies
  decision: AND

data_at_rest:
  layer1:
    aead: aes-256-gcm
    kms: KMS-A
  layer2:
    aead: chacha20-poly1305
    kms: KMS-B

attestation:
  require:
    - sigstore-fulcio
    - org-ca
```

---

## 7) Threat Model & Common-Mode Kill Switches

**Adversaries:** APT with supply-chain access; insider with privileged creds; cryptographic downgrade attacks; cloud control-plane compromise.

**Common-Mode Risks & Mitigations:** - *Shared library vuln* (e.g., *OpenSSL*) → Use different TLS stacks (BoringSSL vs rustls), different AEADs. - *Single CA/KMS breach* → Split vendors/roots; rotate quorum. - *Cloud provider control-plane bug* → Cross-account + cross-region + shadow control plane. - *Policy engine bug* → Dual consensus with canary deny on anomaly.

**Kill Switches:** - Per-layer hard failure defaults to *deny* for privileged ops - Emergency policy: force AND→OR only via out-of-band M-of-N key ceremony with board-level approval

---

## 8) Verification & Testing

- **Tabletop & Red Team:** Simulate single-layer failures; verify others hold.
  - **Fault Injection:** Disable one HSM or policy engine; ensure system continues with reduced capability.
  - **Synthetic Compromise Runs:** Introduce signed but invalid artifact to test twin attestation.
  - **Metrics:** see §9.
- 

## 9) Metrics & SLOs

- **Independence Score (IS):**  $\geq 0.85$  pre-prod;  $\geq 0.9$  for Tier-1 systems.
  - **Common-Mode Risk Index (CMRI):** measured via shared-dep graph; target  $\leq 0.1$ .
  - **Attack Path Reduction (APR):**  $\geq 70\%$  fewer valid single-path exploits vs baseline.
  - **Dual Decision Coverage:**  $\geq 99\%$  of privileged actions gated by DPC.
  - **Key Diversity Index:**  $\geq 2$  vendors,  $\geq 2$  RNG classes,  $\geq 2$  crypto families in use.
- 

## 10) Rollout Plan

**Phase 0 (2–3 weeks):** Dependency graphing, IS/CMRI baseline, choose dual primitives, pick vendors.

**Phase 1 (4–6 weeks):** Implement P1–P3 for a critical service; wire metrics; run failure drills.

**Phase 2 (6–10 weeks):** Extend to transport (P4) and twin attestation (P5); introduce automated audits.

**Phase 3 (Quarterly):** Expand to all Tier-1 services; enforce org policy: any new high-risk control must be coprime with existing stack.

---

## 11) Governance & Audit

- Quarterly **Independence Review** with sign-off on IS/CMRI.
  - SBOM and provenance checks from *two* independent verifiers.
  - External anchoring of logs; third-party attest review.
- 

## 12) Limitations & Tradeoffs

- Complexity and latency overhead; dual encryption and consensus add cost.
- Operational burden: two stacks to maintain; require clear runbooks.
- Some domains can't easily achieve runtime independence (e.g., single cloud kernel).

Mitigate via scoping (Tiered application), automation, and clear exemption process with compensating controls.

---

## 13) Visuals (Placeholders)

- **Layered Stack Diagram:** user→edge→app→data with dual paths highlighted.
  - **Independence Matrix:** rows=layers, cols=dimensions (crypto/trust/code/runtime/ops/supply/control), heatmap of independence.
- 

## 14) Appendix: Coprime Analogy

- In number theory, if factors are **coprime**, divisibility properties compose cleanly: to defeat the product you must defeat each prime factor separately.
- In security, we emulate this by engineering **pairwise-independent failure modes** so that a single exploit cannot transitively defeat all layers.

**Blueprint Outcome:** A repeatable method to design, measure, and audit independence across security layers, turning “defense-in-depth” from a slogan into a quantifiable, provable architecture.